

① “Mental space continuation” of reinforcement learning

In conventional RL, the **environment** is physically observable. I propose to extend it to the **internal** mental space.

From the traditional RL perspective: one reaches out for an apple, the apple is the **reward**, moving the arm is an **action**. These are all observable:



The foundation of RL is the **Bellman equation**. It can be viewed as a **re-
cursive** formula:


Current state

$$V(x_0) = \max_a \{R + \gamma V(x_1)\}$$

Next state

(2)

It propagates the final state’s reward back to the previous state, and the state before that... just like in a chess game... and so on until the very first move□




In other words, the reward of getting an apple, back-propagates to the ac-
tion of reaching out an arm for the apple. So far so good. And we continue
this process back to the **chain of thoughts** that decided to reach for an
apple:

□□□

$$\rightarrow \square\square\square\square\square \rightarrow \square\square\square\square\square\square\square \rightarrow \square\square\square\square \rightarrow$$

(4)

In other words, we turn our **internal** mental states “inside-out”, viewing
them as **external** states, on an equal footing:



And this is exactly analogous to the propagation of rewards in a chess
game. In other words, we can apply techniques of RL to learn the contents
of mental space, thus providing a very rigorous foundation for AGI.

②

dated through **Hebbian learning**



RL 0000 0000 20



² 0000000000000000 RL 000000 00000000000000 000000 00000/000

□□□□□□□□□□ (□□□) □□□□□
□□□□□□□□ □□□□□



(8)

³□□ $T \models M$ □□□□ M is a model of T ; T is a theory of M . □□ □□□□ □□□□□□□□□□